Documentation of Tools for Noise Removal from Pyrosequenced Amplicons (AmpliconNoiseV1.26)

30 January 2011

Usage and reference:

AmpliconNoise is a collection of programs for the removal of noise from 454 sequenced PCR amplicons. Typical usage involves three steps:

1) the removal of noise from the sequencing itself,

2) the removal of PCR point errors, and

3) removal of chimeric sequences using the program Perseus.

While steps 2-3 can be used on amplicon reads generated using most sequencing technologies, step 1 is currently only supported for sequences generated using 454 pyrosequencing.

When presenting work using AmpliconNoise, please refer to the following citation:

Quince *et al* (2011), 'Removing noise from pyrosequenced amplicons.', *BMC Bioinformatics* **12**, 38.

Installation:

Requirements

The programs have been tested on MacOsX and Linux – Windows is not supported. A cluster is not necessary but reasonable size data sets will only run on a cluster or good server. A version of Message Passing Interface (MPI) is necessary to install the programs. Open MPI is a good choice:

http://www.open-mpi.org/

In addition, the chimera checker Perseus requires that both MAFFT and the Gnu Science Library are installed:

http://mafft.cbrc.jp/alignment/software/

http://www.gnu.org/software/gsl/

The proprietary program *sffinfo* from 454 Genomics (Roche) is optional, but the convenience scripts *RunTitanium, RunFLX, RunPreSplit* and *RunFLXPresplit* assume that this is installed. *sffinfo* is part of the analysis software and drivers that are shipped with the sequencing machines from Roche and can possibly be obtained from your sequencing facility, or ask them to deliver plain text versions of all flowgram (SFF) files. Alternatively, you can use the free software Flower by Ketil Malde (http://blog.malde.org/index.php/2009/07/03/a-set-of-tools-for-working-with-454-sequences) or the script *process_sff.py* of QIIME (http://www.qiime.org/) to convert the SFF-files to plain text versions. The convenience scripts are still run referring to the name of the SFF-file but will skip the initial parsing step. If using Flower, the script *RunTitanium.sh* needs to be edited by changing "SplitKeys.pl" to "SplitKeysFlower.pl" (line 52).

Installation procedure

First unzip the programs:

```
unzip AmpliconNoiseV1.26.zip
```

To compile the programs, move into the top directory and type:

```
make clean
make
```

Any errors here may require changing the default C- (*cc*) and C-MPI compilers (*mpicc*) in the individual makefiles associated with the executables.

If the programs compile without errors, type:

```
make install
```

This will place the executables in the bin directory. This directory and the Scripts directory need to be added to your $PATH in order to run the programs from the command line. If you unzip *AmpliconNoiseV1.23.zip* in your home directory ($HOME) then this command should be added to your *.bashrc* or *.profile* or equivalent. Edit this file and add these two line:

```
export PATH=$HOME/AmpliconNoiseV1.26/bin:$PATH
export PATH=$HOME/AmpliconNoiseV1.26/Scripts:$PATH
```

You should also set environment variables to specify the location of look-up tables used by the programs. These define the noise distributions. The following

commands ensure that the file *LookUp_Titanium.dat*[*] is always used for *PyroDist* and *PyroNoise* and *Tran.dat* by *SeqDist* and *SeqNoise*. Having set these the programs can be run anywhere otherwise they can only be run from inside the bin directory. Add the two following lines to your *.bashrc* or *.profile* located in your home directory:

```
export AMPLICON_NOISE_HOME=$HOME/AmpliconNoiseV1.26/
export
PYRO_LOOKUP_FILE=$HOME/AmpliconNoiseV1.26/Data/LookUp_Ti
tanium.dat
export
SEQ_LOOKUP_FILE=$HOME/AmpliconNoiseV1.26/Data/Tran.dat
```

Then either open a new terminal window or source the file to activate these environment variables:

```
source ~/.bashrc
```

*If working with sequencing data generated using an older Pyrosequencing protocol or machine such as (non-Titanium) GS FLX or GS 20, then the *file LookUp_E123.dat* located in the same directory should be used instead of *LookUp_Titanium.dat*

Testing the installation by running an example analysis

The directory Test contains the shell script *Run.sh* which will run through the entire de-noising process for a single dat file. A smallish file consisting of 2,094 GS FLX reads, which will process on a reasonably new MacBook in ten or twenty minutes *C005.dat* is included. This should be run as follows:

```
./Run.sh C005.dat
```

If this works correctly the de-noised file *C005_s60_c01_T220_s30_c08_cd.fa* with just 18 sequences will be generated. The file *C005_s60_c01_T220_s30_c08_cd.mapping* will map these back to the original reads. Other files reflecting the intermediate steps are also generated but in general they can be ignored. The list file giving complete linkage OTUs for these sequences is also produced *C005_s60_c01_T220_s30_c08.list.*

Running AmpliconNoise on medium-size datasets
In the directory *Scripts* in the AmpliconNoise installation directory, there are a number of scripts for running the typical analysis workflow.

The scripts *RunTitanium.sh* and *RunPreSplit.sh* are found in the directory *Scripts*. These can be executed directly on the raw output from pyrosequencing, which is

supplied as flowgram (SFF) files (file suffix ".sff"). These scripts will run in a reasonable time for datasets of with to about 30,000 reads sharing the same (or no) barcode, but this also depends on the evenness of the dataset and the memory available. For larger datasets, see the section "Running Larger Datasets" below.

RunTitanium.sh and RunFLX.sh

These script are used when if several barcoded sequence datasets were provided by the sequencing facility in the same flowgram (SFF) file. It is also used for medium-sized datasets not using barcoded primers. *RunTitanium.sh* is used for GS FLX Titanium or GS FLX Titanium+ datasets whereas *RunFLX.sh* is used for GS FLX or GS20 datasets.

In addition to the flowgram file, a comma-separated text-file named *keys.csv* needs to be present in the same directory. This file should contain one line for each barcode containing the sample name, comma and the barcode sequence, excluding the primer sequence used, e.g.:

```
Sample07,CTCGCGTGTC
```

If no such file is supplied, it is assumed that the sequence data contains no barcodes. In addition, a primer sequence has to be supplied, either by passing it as a second argument to the script or by making a FASTA-formatted file containing only the primer sequence named *primer.fasta*, e.g.:

```
>787F
ATTAGATACCCNGGTAG
```

Note that this file may contain degenerated base characters, such as 'N'.

As a default, the script is run using 4 parallel processes. To change this value, edit the script or make a local copy of it and edit that. Change line "nodes=4" to as many processes as you would like to run. Larger datasets can be processed faster with this script on a more powerful computer or cluster by increasing this value to e.g. 32 or more.

The script is then executed simply by supplying the name of the flowgram file, e.g.

```
RunTitanium.sh MySamples.sff
```

RunPreSplit.sh and RunFLXPreSplit.sh

These scripts are used for barcoded sequence datasets that have already been split up by the sequencing facility, so that one flowgram (SFF) file contains one sample. They are used like the *RunTitanium.sh* and *RunFLX.sh* scripts described above, except that no *keys.csv* file is needed.

Output of AmpliconNoise

The different processes of the AmpliconNoise workflow generate several files, most important of which are (using "SampleX" as an example sample name):

SampleX_F_Good.fa
This file contains the unique sequences after removal of sequencing noise, PCR point errors and chimeras, in FASTA format. The last number of each sequence name given in the FASTA header, indicated after the underscore character, represents the number of reads that are most likely to share this unique sequence is after cleaning. For example the fasta header ">LA_RNA_s60_c01_T400_P_BC_s30_c08_0_8" indicates that this sequence represent eight reads, from the sample *LA_RNA*.

SampleX_OTUs_0.03.fasta
This file contains representative sequences for each OTU (Operational Taxanomic Unit) after a 3% maximum linkage clustering of the unique, de-noised sequences. The number of reads is indicated just like in *SampleX_F_Good.fa*

AN_stats.txt
This tab-separated text file contains statistics about the run, such as number of reads and unique sequences before and after de-noising and chimera filtering, number of OTUs and two diversity estimates (the Shannon diversity index / entropy and Simpson's diversity index given as 1-D).

SampleX_[..]_cd.mapping
This file maps each unique sequence back to the name of its reads.

SampleX.raw.fasta and SampleX.raw.fasta.qual
Generated by *RunPreSplit.sh* only, these contain the sequences and base qualities of the reads before initial quality filtering and de-noising.

Step-by-step description of a typical workflow

This section explains the script RunTitanium.sh step by step. The other scripts (*RunFLX*, *RunPreSplit* and *RunFLXPreSplit*) follow the same general outline. Differences and alternatives are described briefly.

```
export bc=keys.csv
export nodes=4
export otu_dist=0.03
```

These values specify the name of the barcodes metadata file *keys.csv*, the number of nodes to be used and the maximum sequence distance for maximum linkage clustering, after de-noising.

After this, the workflow is started by reading the primer sequence (unless given as an argument to the script). A plain-text version of the flowgram (SFF) file is then generated unless such a file already exists:

```
if [ ! -f ${stub}.sff.txt ]; then
    echo "Generating .sff.txt file"
    sffinfo $1 >${stub}.sff.txt
fi
```

The plain-text flowgram file is then quality filtered and parsed into one or more *.dat* files containing only the identifiers and flow values of those reads that pass quality filtering. The exact procedure depends on which script is used. If a barcode-file is supplied, *RunTitanium* and *RunFLX* will split the dataset into one *.dat* file for each barcode. Only exact matches to the given barcodes are retrieved. If *keys.csv* is not present or if using *RunPreSplit* or *RunFLXPreSplit,* the barcode sequences are not checked, but the length of the barcode is saved in the last line of the *.stat.txt* file.

The pre-filtering control also removes all reads with fewer than 360 flows before the first empty flow cycle or degenerate base (flow intensity between 0.5 and 0.7). Pre-filtering is either carried out by the java-class *ampliconflow.sff.FlowsFlex,* or the perl-scripts *FlowsMinMax.pl*, *FlowsFA360.pl*, *CleanMinMax.pl* or *Clean360.pl* depending on the script and dataset type. FLX reads are cropped at 360 flows and Titanium at 720.

Following pre-filtering, all steps are repeated for each sample using *RunTitanium / RunFLX.* First, distances between flowgrams are calculated using *PyroDist*:

```
echo "Running PyroDist for ${stub}"
mpirun -np $nodes PyroDist -in $file -out ${stub} >
${stub}.fout
```

Then, hierarchical clustering with complete linkage is carried out using *FCluster* to provide input file for *PyroNoise.* (Some intermediate files are also removed):

```
echo "Clustering .fdist file"
FCluster -in ${stub}.fdist -out ${stub}_X > ${stub}.fout

rm ${stub}.fdist
rm ${stub}_X.otu ${stub}_X.tree
```

Next, the flowgrams are iteratively clustered according to the EM algorithm implemented in *PyroNoise*, to remove pyrosequencing noise. By default, an initial clustering cut-off of 0.01 and cluster size of 60.0 is used (see Section Programs for details).

```
echo "Running PyroNoiseM"
```

```
mpirun $mpiextra -np $nodes PyroNoiseM -din ${stub}.dat
-out ${stub}_s60_c01 -lin ${stub}_X.list -s 60.0 -c 0.01
> ${stub}_s60_c01.pout
```

The ends of reads are often noisy, so next, we truncate these to 400 bp (220 for
FLX reads). This position can be moved to change the balance between
remaining noise and sequence length.

```
Truncate.pl 400 < ${stub}_s60_c01_cd.fa >
${stub}_s60_c01_T400.fa
```

Next, barcodes and primer sequences are removed using the script *cropF.py*.

```
cropF.py  ${stub}_s60_c01_T400.fa $cropFL >
${stub}_s60_c01_T400_P_BC.fa
```

Now we calculate the PCR-error-corrected distances between sequences using
SeqDist:

```
echo "Running SeqDist"

mpirun -np $nodes SeqDist -in ${stub}_s60_c01_T220.fa >
${stub}_s60_c01_T220.seqdist
```

Complete linkage clustering (again using *FCluster*) is carried out to provide input
to *SeqNoise*:

```
echo "Clustering SeqDist output for ${stub}"

FCluster -in ${stub}_s60_c01_T220.seqdist -out
${stub}_s60_c01_T220_S > ${stub}_s60_c01_T220.fcout
```

SeqNoise implements the sequence clustering algorithm that removes PCR
errors:

```
echo "Running SeqNoise"

mpirun -np $nodes SeqNoise -in ${stub}_s60_c01_T220.fa -
din ${stub}_s60_c01_T220.seqdist -lin
${stub}_s60_c01_T220_S.list -out
${stub}_s60_c01_T220_s30_c08 -s 30.0 -c 0.08 -min
${stub}_s60_c01.mapping > ${stub}_s60_c01_T220.snout

rm ${stub}_s60_c01_T220_S.otu
${stub}_s60_c01_T220_S.tree ${stub}_s60_c01_T220.seqdist

ln -s ${stub}_s60_c01_T400_P_BC_s30_c08_cd.fa
${stub}_F.fa
```

Perseus is then used to screen for and remove chimeric sequences. Depending on the targeted sequence region or gene, parameters can be changed and *PerseusD* can alternatively be used, which reduces the false positive rate at the cost of slightly reduced sensitivity.

```
echo "Running Perseus for ${stub}"

Perseus -sin ${stub}_F.fa > ${stub}_F.per

Class.pl ${stub}_F.per -7.5 0.5 > ${stub}_F.class

FilterGoodClass.pl ${stub}_F.fa ${stub}_F.class 0.5
1>${stub}_F_Chi.fa 2>${stub}_F_Good.fa
```

Finally we build OTUs from the de-noised sequences:

```
echo "Clustering OTUs"
mpirun -np $nodes NDist -i -in
${stub}_s60_c01_T220_s30_c08_cd.fa >
${stub}_s60_c01_T220_s30_c08.ndist

FCluster -i -in ${stub}_s60_c01_T220_s30_c08.ndist -out
${stub}_s60_c01_T220_s30_c08 >
${stub}_s60_c01_T220_s30_c08.fcout

echo "Writing otu representatives and statistics"

java amliconflow.otu.OTUUtils -in ${stub}_F_Good.list -
dist $otu_dist -repseq ${stub}_F_Good.fa >
${stub}_OTUs_${otu_dist}.fasta
```

The remaining steps calculate and write statistics and diversity estimates to the file *AN_stat.txt*.

Running larger datasets

The directory *TestFull* contains a shell script that illustrates the de-noising process for a larger sample that needs to be split to allow de-noising. This should only be run on a cluster or good server. It is assumed that a single sample without barcodes is used. The script takes an .sff file as an argument but in case *sffinfo* (a proprietary program from 454 / Roche) is absent we have provided the plain-text version *ArtificialGSFLX.sff.txt*.

The script should be run in general as:

```
./Run.sh My.sff primer
```

..and for test purposes:

```
./Run.sh ArtificialGSFLX.sff
```

For files that have already been split up into one .sff-file per barcode, the script *RunPreSplitXL.sh* is provided (in the *Scripts* directory). This is used just like *RunPreSplit.sh*

Step-by-step description of TestFull/Run.sh

```
#!/bin/bash

defaultPrimer="ATTAGATACCC\w{1}GGTAG" #default primer
nodes=8                               #no. of cluster
nodes to use

sfffile=$1; #first argument name of sff file (necessary)
primer=${2:-$defaultPrimer} #second argument primer as a
Perl regular expression
```

(the second argument should be your primer else it defaults to our 787F)

```
stub=${sfffile%.sff};

echo $stub $sfffile $primer

# first generate sff text file if necessary
if [ ! -f ${sfffile}.txt ]; then
    echo "Generating .sff.txt file"
    sffinfo $sfffile > ${sfffile}.txt
fi
```

(generates text translation of sff file if necessary):

```
#generate flowgram and fasta files
if [ ! -f ${stub}.dat ]; then
    echo "Generating .dat file"
    FlowsFA360.pl $primer $stub < ${sfffile}.txt
fi
```

(extracts filtered dat and sequence files)

Next, a file of unique sequence is generated, using the *FastaUnique* program:

```
#get unique sequences
if [ ! -f ${stub}_U.fa ]; then
    echo "Getting unique sequences"
    FastaUnique -in ${stub}.fa > ${stub}_U.fa
fi
```

The unique sequences are then used to generate an average linkage tree, based on sequence distance, with the program *NDist*:

```
#use NDist to get sequence distances
if [ ! -f ${stub}_U_I.list ]; then
    echo "Calculating sequence distances"
    mpirun -np $nodes NDist -i -in ${stub}_U.fa >
${stub}_U_I.ndist
fi

#use NDist to get sequence distances
if [ ! -f ${stub}_U_I.list ]; then
    echo "Cluster sequences..";
#cluster sequences using average linkage and sequence
weights
    FCluster -a -w -in ${stub}_U_I.ndist -out
${stub}_U_I > ${stub}_U_I.fcout
fi

rm ${stub}_U_I.ndist
```

The average linkage clustering is then then used to split up the .dat file. The parameters –s and –m can be modified in order to change the cluster size (see Scripts below for details):

```
SplitClusterEven -din ${stub}.dat -min ${stub}.map -tin
${stub}_U_I.tree -s 5000 -m 1000 > ${stub}_split.stats
```

Now we can start de-noising each .dat file separately beginning by calculating the flowgram distances:

```
echo "Calculating .fdist files"
for c in C*
do
        if [ -d $c ] ; then
                mpirun -np $nodes PyroDist -in
${c}/${c}.dat -out ${c}/${c} > ${c}/${c}.fout
        fi
done
```

..cluster them:

```
echo "Clustering .fdist files"

for c in C*
do
        if [ -d $c ] ; then
                FCluster -in ${c}/${c}.fdist -out
${c}/${c}_X > ${c}/${c}.fout
                rm ${c}/${c}.fdist
```

```
        fi
done
```

...and denoise them to get sequences:

```
echo "Running PyroNoise"
for dir in C*
do
        if [ -d $dir ] ; then
                mpirun -np $nodes PyroNoise -din
${dir}/${dir}.dat -out ${dir}/${dir}_s60_c01 -lin
${dir}/${dir}_X.list -s 60.0 -c 0.01 >
${dir}/${dir}_s60_c01.pout
        fi
done
```

In this case we then concatenate the sequences together and do a single sequence noise removal step. Alternatively, SeqNoise can be run in the the separate directories before bringing them together for a final noise removal step, to speed up the process and save memory:

```
cat C*/C*_s60_c01_cd.fa > All_s60_c01_cd.fa
cat C*/C*_s60_c01.mapping > All_s60_c01.mapping

Truncate.pl 220 < All_s60_c01_cd.fa >
All_s60_c01_T220.fa


echo "Running SeqDist"
mpirun -np $nodes SeqDist -in All_s60_c01_T220.fa >
All_s60_c01_T220.seqdist

FCluster -in All_s60_c01_T220.seqdist -out
All_s60_c01_T220_S > All_s60_c01_T220.fcout

rm All_s60_c01_T220.seqdist
```

Finally we remove PCR error from our sequences:

```
echo "Running SeqNoise"
mpirun -np $nodes SeqNoise -din All_s60_c01_T220.fa -lin
All_s60_c01_T220_S.list -out All_s60_c01_T220_s30_c08 -s
30.0 -c 0.08 -min All_s60_c01.mapping >
All_s60_c01_T220_s30_c08.snout
```

Programs

FCluster:

-in   string   distance input file name
-out string   output file stub
Options:
-r   resolution
-a   average linkage
-w   use weights
-i   read identifiers
-s   scale dist.

This performs a simple hierarchical clustering. It reads a distance file in text format (-in).

The first line in the text file gives the number of entities to be clustered N. This is then optionally followed by N ids if the (-i) flag is set as separate lines. Otherwise the N(N-1)/2 pairwise distances follow as individual lines. The distances d_{ij} are specified in order i = 1…N, j = 1..i.

The program performs complete linkage clustering as default but average linkage can be specified by the (-a) flag. Average linkage accounting for weights is possible with (-a –w) the weights are then take from the ids which must have format

```
 Name1_Weight1
 …
 NameN_WeightN
```

The program produces three output files *stub.list, stub.otu, stub.tree* when stub is specified by (-out):

stub.list has format (similar to *DOTUR*):
```
 d NClusters Cluster1 .. ClusterN
```

where *d* is the distance at which clusters formed. *N* is the number of clusters at this cutoff and then each cluster is specified as a comma separated list of entries either indexed 0 to N -1 or by ids if the (-i) flag is specified.

stub.otu simply gives the cluster sizes in the same format. Clusters are outputted at separations of 0.01 by default but this can be change by (–r) flag.

stub.tree is the hierarchical in newick tree format

Finally the distances can be scaled by their maximum using the (-s) flag.

Examples:

To perform complete linkage hierarchical clustering:

```
FCluster -in test.fdist -out test_M
```

Or to use average linkage with weights and ids in output:

```
FCluster -i -a -w -in test.ndist -out test_A
```

(this requires distance file with ids)

FClusterM:

-in   string   distance input file name
-out string   output file stub
Options:
-r   resolution
-a   average linkage
-w   use weights
-i   read identifiers
-s   scale dist.

This performs a simple hierarchical clustering. It reads a distance file in text format (-in) that has a full distance matrix. The first line in the text file gives the number of entities to be clustered N. This is then optionally followed by N ids if the (-i) flag is set as separate lines. Otherwise the N*N pairwise distances follow as individual lines. The distances d_{ij} are specified in order i = 1...N, j = 1..N. For clustering this matrix is converted into its symmetric equivalent d_ij = 0.5*(d_ij + d_ji). This is suitable for clustering the output of *SeqDistM*.

FClusterN:

-in   string   distance input file name
-out string   output file stub
Options:
-r   resolution
-a   average linkage
-w   use weights
-i   read identifiers
-s   scale dist.

This is simply a more efficient version of FCluster due to David Hunt at Tessella funded through a Unilever project. It produces identical results to FCluster but much faster. In the next update it will replace the current FCluster code.

FastaUnique:
-in   string   input file name

This program simply dereplicates a fasta file of sequences. Sequences of different length are only compared up to the smaller length and if identical up to that smaller length are judged the same sequence. Dereplicated sequences with ids that are a combination of the founding sequence id and the number of identical sequences found i.e.

```
>founderID_weight
```

The mapping of sequences to the uniques is given by a .map file generated with the name *fastaname.map* where *fastaname* is obtained by parsing .fa of the original file name. This has a line for each unique sequence in format:

```
OriginallIdx, NewIdx, ParentID, I:
Idx_1,…Idx_I:ID_1,…,ID_I
```

..where I is the number of sequences mapping to the unique.

Example:

```
FastaUnique -in Test.fa > Test_U.fa
```


NDist (pairwise Needleman-Wunsch sequence distance matrix from a fasta file)

-in   string   fata file name
Options:
-i output identifiers

This program generates a distance matrix from a fasta file of the format required by *FCluster*. It uses a simple implementation of the exact Needleman-Wunsch algorithm to perform pairwise alignments using a fixed gap penalty of 1.5. Distances are then calculated according to the '*QuickDist'* algorithm basically counting mismatched nucleotides as a distance of one and with a cost of one for a gap regardless of length and then normalizing by number of comparisons (Huse *et al.* Genome Biology 2007). Output is to standard out.

The only option (-i) is to output identifiers suitable for running *FCluster* with –i.

This is an MPI program allowing the calculation of distances to spread across multiple cores and/or nodes.

Example:

```
mpirun –np 32 NDist –in Test.fa > Test.ndist
```

Perseus (slays monsters)

-sin   string   seq file name
Options:
-s   integer
-tin   string   reference sequence file
-a   output alignments
-d   use imbalance
-rin   string   lookup file name

The Perseus algorithm given an input fasta file (-sin) takes each sequence in turn and searches for the closest chimeric match using the other sequences as possible parents. It finds the optimum parents and breakpoints. It only searches for parents amongst species of equal or greater abundance where abundance is obtained from the fasta ids:

```
>ID_weight
```

Never run multiple copies of *Perseus* in the same directory! The (-a) flag outputs all the chimeric alignments and is useful for verifying if sequence truly is chimeric. The (-d) flag uses a slightly different algorithm including a penalty for imbalance on branches of the tree formed by the chimera and parents which may give better results in some instances. Perseus uses a nucleotide transition file and (–rin) allows this file to be set otherwise it defaults to the *SEQ_LOOKUP_FILE* variable and if this is not set the header variable *LOOKUP_FILE* which is set to "*../Data/Tran.dat*".

We recommend removing degenerate primers before running *Perseus*.

It produces a lot of info but ... the critical portion are the x=12th, y=13th, and z=14th tokens. If x < 0.15 and y >= 0.0 and z is larger than about 15 then this is a chimera.

The (-s) controls skew i.e. how much greater in frequency a sequence has to be to be a putative parent. This default to one – higher values can reduce the false positive rate.

The (-tin) option allows sequences other than the queries to be used as references. This can be used to split a file for running across threads or on a cluster (see example below).

Example usage:

```
sed 's/^ATTAGATACCC\w{1}GGTAG//'
C005_s60_c01_T220_s30_c08_cd.fa >
C005_s60_c01_T220_s30_c08_P.fa

Perseus -sin C005_s60_c01_T220_s30_c08_P.fa >
C005_s60_c01_T220_s30_c08_P.per
```

To split a fasta file into four sections each in its own directory and then run *Perseus* in the background on each separately before recombining the output:

```
Split.pl Uneven1_s25_P.fa 4

cd Split0
Perseus -sin Split0.fa -tin ../Uneven1_s25_P.fa >
Split0.per&

cd ../Split1
Perseus -sin Split1.fa -tin ../Uneven1_s25_P.fa >
Split1.per&

cd ../Split2
Perseus -sin Split2.fa -tin ../Uneven1_s25_P.fa >
Split2.per&

cd ../Split3
Perseus -sin Split3.fa -tin ../Uneven1_s25_P.fa >
Split3.per&

../Scripts/Join.pl Split*/*per > Uneven1_s25_P.per
```

To classify sequences use Class.pl with suggested parameters for V5:

```
Class.pl C005_s60_c01_T220_s30_c08_P.per -6.6925 0.5652
> C005_s60_c01_T220_s30_c08_P.class
```

..this generates a file with format:

```
seqname x y z probability_of_being_chimeric
```

We can split up the original fasta file at 50% probability of being chimeric:

```
FilterGoodClass.pl C005_s60_c01_T220_s30_c08_P.fa
C005_s60_c01_T220_s30_c08_P.class 0.5 2>
C005_s60_c01_T220_s30_c08_Good.fa >
C005_s60_c01_T220_s30_c08_Chi.fa
```

PerseusD (slays monsters)

-sin **string** **seq file name**
Options:
-c **float,float** **set alpha,beta default = -5.54,0.33**
-s **integer** **set skew default = 2**
-tin **string** **reference sequence file**
-a **output alignments**
-b **do not use imbalance**
-rin **string** **lookup file name**

PerseusD differs in algorithm and output from *Perseus*. It only tests against parents that have been classified as non-chimeric. It also only tests for possible parents amongst sequences that are at least twice as abundant as the query. These changes reduce false positives but at the cost that sensitivity is also slightly reduced. They were inspired by the strategy adopted in *uchime* (Edgar *et al.* 2011 'UCHIME improves sensitivity and speed of chimera detection', *Bioinformatics*). This program should be preferred when a few chimeras can be tolerated and false positives cannot. Unlike *Perseus* it needs to perform classification itself. Usage is just like *Perseus* except that it generates .class-files rather than .per equivalent to running *Perseus* and then *Class.pl*:

Example usage:

```
Perseus -sin C005_s60_c01_T220_s30_c08_P.fa >
C005_s60_c01_T220_s30_c08_P.class
```

The out format is therefore of this form:

```
SeqName x y z p
```

..where *p* is the probability of the sequence being chimeric. Never run multiple copies of *PerseusD* in the same directory! *PerseusD* uses the imbalance penalty as default. The (-b) flag turns this off. The flag (-c alpha,beta) allows different alpha and beta parameters to be passed to the program these default to values for the V5 region trained through logistic regression. These work well generally though. Other parameters are as for *Perseus*.

PyroDist (pairwise distance matrix from flowgrams)

-in	string	flow file name
-out	stub	out file stub
Options:		
-ni		no index in dat file
-rin	string	lookup file name

This program calculates a distance matrix between flowgrams. Input (-in) is to a .dat file containing flowgrams in a simple format. The first line has the number of flowgrams followed by the number of flows: N M. Each of the N flowgram entries has the format: id length1 flow1 flow2 ... flowM where id is just an identifier, length is the number of 'clean' flows, followed by all M flows (although only length will ever be used).

The distances are calculated according to the algorithm in Quince *et al.* 2009 except that alignment of flowgrams no longer occurs. This requires a look-up table for the intensity distributions about the homopolymer length. By default this is read in from a file set in the header file by the constant LOOKUP_FILE which is set to "../Data/LookUp_E123.dat" a well configured distrubution for 454

GSFLX implementation. Consequently the program can only be run from the bin directory to maintain this relative path. However, to allow the program to run anywhere the environment variable *PYRO_LOOKUP_FILE* can be set as described in the installation instructions or the path to a lookup file can be passed with the (-rin) flag.

The optional flag (-ni) is necessary if the flowgram file contains no ids.

Output is to a distance matrix in flat format of name stub.fdist where stub is set by the (-out) flag. Status information is sent to standard out and this can be safely ignored if the program runs correctly.

This is an MPI program allowing the calculation of distances to spread across multiple cores and/or nodes.

Example:

```
mpirun –np 32 PyroDist –in Test.fa –out Test >
Test.fdout
```

This generates the distance matrix *Test.fdist*

PyroNoise (clusters flowgrams without alignments)
-din	**string**	**flow file name**
-out	**string**	**cluster input file name**
-lin	**string**	**list file**
Options:		
-v	**verbose**	
-c	**double**	**initial cut-off**
-ni		**no index in dat files**
-s	**double**	**precision**
-rin	**file**	**lookup file name**

This program uses an EM algorithm to construct de-noised sequences by clustering flowgrams as described in Quince et al. 2009 but without alignments. It takes as input (-din) a flowgram file of the format described above and an initial hierarchical clustering (-lin) generated by running *FCluster* on the output of *PyroDist*. Output files are generated with the stub specified by flag (-out).

The cut-off for the initial clustering is specified by (-c) generally this should be quite small 0.01 is a good value for most data sets. The paramter (-s) controls the cluster size. The larger this is the tighter the clusters – 60.0 is a reasonable value here but smaller may remove more pyrosequencing noise. If these parameters are not set they default to these values.

The parameter (-rin) allows a look up file to be specified otherwise the program uses the environment variable PYRO_LOOKUP_FILE if that is not set it defaults to the global variable LOOKUP_FILE found in PyroNoise.h currently

"../Data/LookUp_E123.dat". This will work provided the executable is run from the bin directory to maintain this relative path to the files in ../Data.

The option (-v) outputs extra debug information to standard out.

Information on cluster convergence is output to standard out and after running the program produces a number of files:

1) *stub_cd.fa*: a fasta file of de-noised sequences. The ids are formed as "*>stub_index_weight*" where weight are the number of reads mapping to that sequence, and index is just an arbitrary cluster number.

2) *stub_cd.qual:* qualities for the denoised sequences see Quince et al. (unpublished).

3) *stub.mapping*: contains a line for each de-noised sequence giving the read that characterizes that sequence followed by a tab separated list of flowgram reads (specified by their ids read from dat file) that map to it.

4) directory stub: contains a fasta file for each de-noised sequence, *i_index.fa*, of reads that map to it.

This is an MPI program allowing the calculation of distances to spread across multiple cores and/or nodes.

Example:

```
mpirun –np 32 PyroNoise -din Test.dat -out Test_s60_c01
-lin Test_X.list -s 60.0 -c 0.01 > Test_s60_c01.pout
```

PyroNoiseM

This version of *PyroNoise* has the exact same usage as above but stores flowgram distances in memory. It is useful for Titanium data where the calculation of these distances may be the limiting step.

SeqDist (pairwise distance matrix from a fasta file)

-in   string   fasta file name
Options:
-i output identifiers
-rin   string   lookup file name

This program generates a distance matrix of the format required by *FCluster* from a fasta file. It uses a an implementation of the exact Needleman-Wunsch algorithm to perform pairwise alignments. Distances account for nucleotide transition probabilities as a result of PCR errors. There is a different cost for homopolymer (4.0) and normal gaps (15.0). The probabilities, actually –log of, are read from a look up table. By default this is from a file set in the header file by the constant *LOOKUP_FILE* which is set to "*../Data/Tran.dat*" configured for a standard polymerase. Consequently the program can only be run from the bin directory to maintain this relative path. However, to allow the program to run anywhere the environment variable *SEQ_LOOKUP_FILE* can be set as described in

the installation instructions or the path to a lookup file can be passed with the (-rin) flag.

The option (-i) is to output identifiers suitable for running *FCluster* with –i.

This is an MPI program allowing the calculation of distances to spread across multiple cores and/or nodes.

Example:

```
mpirun –np 32 SeqDist –in Test.fa > Test.seqdist
```

SeqDistM

This version of *SeqNoise* has the exact same usage as above but generates an asymmetric distance matrix NXN distance matrix that is appropriate for *SeqNoiseM*.

SeqNoise (*clusters sequences)*

-in   string   fasta sequence file name
-din   string   sequence distances file name
-out   string   cluster input file name
-lin   string   list file
Options:
-min   mapping file
-v   verbose
-c   double   initial cut-off
-s   double   precision
-rin   string   lookup file name

This program uses an EM algorithm to remove PCR noise by clustering sequences as described in Quince et al. (2011). The same distance metric as described in *SeqDist* is used. It takes as input (-in) a fasta file (with frequencies defined in ids as *>id_weight*), (-din) a flat matrix of sequence distances generated by *SeqDist* and an initial hierarchical clustering (-lin) generated by running *FCluster* on the output of *SeqDist*. Output files are generated with the stub specified by flag (-out).

The cut-off for the initial clustering is specified by (-c) generally this should be quite large 0.08 is a good value for most data sets. The paramter (-s) controls the cluster size. The larger this is the tighter the clusters – 30.0 is a reasonable value here but smaller may remove more noise and larger allow high resolutions OTUs. If these parameters are not set they default to these values.

The parameter (-rin) allows a look up file to be specified otherwise the program uses the environment variable *SEQ_LOOKUP_FILE* if that is not set it defaults to

the global variable LOOKUP_FILE found in *SeqNoise.h* currently *"../Data/Tran.dat".* This will work provided the executable is run from the bin directory to maintain this relative path to the files in *../Data.*

The option (-v) outputs extra debug information to standard out.

The option (-min) allows a mapping file from a previous *PyroDist* step to be input. If used the program will use this information to map denoised sequences back to the original flowgram ids.

Information on cluster convergence is output to standard out and after running the program produces a number of files:
1) *stub_cd.fa*: a fasta file of de-noised sequences. The ids are formed as *">stub_index_weight"* where weight are the number of sequences mapping to that sequence, and index is just an arbitrary cluster number.
2) *stub.mapping*: contains a line for each de-noised sequence giving the input sequence defining the denoised cluster followed by a tab separated list of input sequences that map to that sequence.
3) directory *stub*: contains a fasta file for each de-noised sequence, *i_index.fa*, of sequences that map to it.
4) Optional on (-min) if a mapping file is input then a file *stub_cd.mapping* containing a line for each de-noised sequence giving the id followed by a tab separated list of original reads that map to it.

This is an MPI program allowing the calculation of distances to spread across multiple cores and/or nodes.

Example:

```
mpirun -np 32 SeqNoise -in Test_s60_c01_T220.fa -din
Test_s60_c01_T220.seqdist -lin Test_s60_c01_T220_S.list
-out Test_s60_c01_T220_s30_c08 -s 30.0 -c 0.08 -min
Test_s60_c01.mapping > Test_s60_c01_T220.snout
```

SeqNoiseM

This version of *SeqNoise* has the exact same usage as above but uses a slightly different algorithm for the centroid construction which will prefer longer sequences for centroid clusters. It may be preferred for Titanium data if read lengths are very uneven (std dev > 100) it requires input from *SeqDistM*.

SplitClusterEven

-din   string dat filename
-min   string map filename
-tin   string tree filename
-s   split size
-m   min size

This program splits up *dat* files (-din) using a tree generated on unique sequences (-tin) input as a .tree file. The mapping of unique sequences to reads in the dat file is specified by a .map file (-min). The tree is the split in such a way as to maintain a maximum (-s) and minimum (-m) cluster size (measured on unique reads). The parameters *–s 2500* and *–m 250* will likely produce *dat* files of a good size although you should play around with these. The dat files are placed in directories labeled C000, ..,C00N+ where N is the number of clusters and the + simply indicates that this will be an aggregation of all small clusters.

Scripts:

Some useful Perl scripts are also provided in the Scripts directory:

FlowsFA.pl

This extracts flowgrams from a plain-text flowgram (*.sff.txt*) file. It takes the primer as a first argument and an output stub as the second. It reads from standard input. It should be used for GS FLX reads. For example:

```
FlowsFA.pl ATTAGATACCC[ACTG]GGTAG ArtificialGSFLX <
ArtificialGSFLX.sff.txt
```

..will generate the filtered *.dat* flowgram file *ArtificialGSFLX.dat* and a fasta file of the corresponding sequences *ArtificialGSFLX.fa.* Filtering requires that a minimum sequence length of 204 (changed by altering variable $minLength) including key and primer is achieved before the first noisy signal (0.5-0.7 or no signal across all four bases). Flowgrams are then truncated at this point. If keys are used simply pass the entire key–linker–primer sequence to this script or use *SplitKeys.pl* described below.

FlowsFA360.pl

This extracts flowgrams from the text translation of a .sff.txt. It takes the primer as a first argument and an output stub as the second. It reads from standard input. It should be used for GS FLX reads. For example:

```
FlowsFA360.pl ATTAGATACCC[ACTG]GGTAG ArtificialGSFLX <
ArtificialGSFLX.sff.txt
```

..will generate the filtered .dat flowgram file *ArtificialGSFLX.dat* and a fasta file of the corresponding sequences *ArtificialGSFLX.fa.* Filtering requires that a minimum flowgram length of 360 including key and primer is achieved before the first noisy signal (0.5-0.7 or no signal across all four bases). All flowgrams are then truncated at 360. If keys are used simply pass the entire key – linker – primer sequence to this script or use *SplitKeys.pl* described below.

FlowsMinMax.pl

This extracts flowgrams from the text translation of a .sff.txt file. It takes the primer as a first argument and an output stub as the second. It reads from standard input. It should be used for Titanium reads. For example:

```
FlowsMinMax.pl    ACACACGTCGACTCCTACGGGAGGCAGCAG
TitaniumV3 < TitaniumV3.sff.txt
```

…will generate the filtered .dat flowgram file TitaniumV3.dat and a fasta file of the corresponding sequences *TitaniumV3.fa* for a key `ACACACGTCG` and primer `ACTCCTACGGGAGGCAGCAG.` Filtering requires that a minimum flowgram length of 360 including key and primer is achieved before the first noisy signal (0.5-0.7 or no signal across all four bases). All flowgrams are then truncated at 720. If keys are used simply pass the entire key – linker –primer sequence to this script in upper case or use *SplitKeys.pl* described below.

CountFasta.pl

Gives total read number mapping to a fasta file with weighted ids. For example:

```
CountFasta.pl < Test_s60_c01_cd.fa
```

Truncate.pl

Truncates sequences in a fasta file e.g.:

```
Truncate.pl 220 < Test_s60_c01_cd.fa >
Test_s60_c01_T220.fa
```

SplitKeys.pl

Separates out an sff file read from standard input according to barcode sequences. Requires a file *keys.csv* with format:

```
SampleName1, Barcode1
…
SampleNameN, BarcodeN
```

The primer is the first argument of the script. The second is the *keys.csv* file. This script generates .raw files that then have to be filtered and reformatted using Clean360.pl. A shell script Clean.sh shows how to do this for multiple raw data files. Reads that do not match to any tag are output to standard error. Any linkers must be included in the barcodes.

```
./SplitKeys.pl TGCTGCCTCCCGTAGGAGT keys.csv <
FV9NWLF01.sff.txt 2> Err.fa
```

SplitKeysFlower.pl

Separates out a flower file generated generated by Ketil Malde's program Flower (http://blog.malde.org/index.php/2009/07/03/a-set-of-tools-for-working-with-454-sequences). It reads from standard input according to barcode sequences. Requires a file *keys.csv* with format:

```
SampleName1, Barcode1
…
SampleNameN, BarcodeN
```

The primer is the first argument of the script. The second is the *keys.csv* file. This script generates .raw files that then have to be filtered and reformatted using *Clean360.p*l. A shell script *Clean.sh* shows how to do this for multiple raw data files. Reads that do not match to any tag are output to standard error. Any linkers must be included in the barcodes:

```
./SplitKeysFlower.pl TGCTGCCTCCCGTAGGAGT Keys.csv <
FV9NWLF01.flower.txt 2> Err.fa
```

Qiime_Typical.pl

Generates OTU consensus sequences with format suitable for *Qiime*. Takes fractional sequence difference for OTU construction as the first argument, a fasta file of denoised sequences for the second and list file from *NDist* for the third. See the tutorial for more information.

Example:

```
./Qiime_Typical.pl 0.03 All_Good.fa All_Good.list >
All_Good_C03_Q.fa
```

Qiime_OTU.pl

Generates Qiime OTU tables. Takes fractional sequence difference for OTU construction as the first argument, RDP taxonomic classifications as second and sample suffix for third. Generate classifications from using Qiime using:

```
assign_taxonomy.py -i All_Good_C03_Q.fa
```

Example:

```
./Qiime_OTU.pl 0.03
rdp_assigned_taxonomy/All_Good_C03_Q_tax_assignments.txt
TS < All_Good.list > All_Good_C03.qiime
```

The file *All_Good_C03.qiime* can now be used directly in Qiime as an OTU table.